

# MidiComb

## A CP-based Approach to Music Generation

---

Francesco Ballerini

Università di Bologna

# Motivation

---

Hyun et al. introduce **ComMU**: a dataset of short musical samples built for the task of **combinatorial music generation**.


## Combinatorial Music Generation

Given a set of short musical **samples** together with specific **metadata** describing their properties, combine them to create a complete piece of music.

If samples share a subset of the metadata — bpm, key, time signature, number of measures, genre, rhythm, and chord progression — they are guaranteed to sound **harmonious** when combined together.

# Combinatorial Music Generation

- **Stage 1:** train a machine learning model on ComMU to generate new samples given the desired metadata
- **Stage 2:** combine the generated samples into a complete piece of music



(a) stage1

Main melody		note sequence(1)	note sequence(4)
Sub melody			
Accompaniment	note sequence(2)	note sequence(2)	note sequence(2)
Bass			
Riff		note sequence(3)	note sequence(3)
Pad	note sequence(5)		

(b) stage2

# Our contribution: MidiComb

Hyun et al. only focus on stage 1: they train a **Transformer-XL** on ComMU and show the quality of the generated samples, leaving stage 2 for future work.

**MidiComb** addresses stage 2 of the combinatorial music generation task by:

1. Taking in input the desired shared metadata values
2. Querying ComMU for a set of samples satisfying those requirements
3. Generating a complete musical composition by combining those samples

# Implementation

---

# Music Generation as a Scheduling Problem

The sample arrangement problem is modeled as a **scheduling problem**, analogous to the **job shop problem**:

- Each **track role** (main melody, sub melody, etc.) represents a **machine**
- Each **sample** represent a **task**

Main melody		note sequence(1)	note sequence(4)
Sub melody			
Accompaniment	note sequence(2)	note sequence(2)	note sequence(2)
Bass			
Riff		note sequence(3)	note sequence(3)
Pad	note sequence(5)		

# Music Generation as a Scheduling Problem

- **Constraints:** *how* samples are combined
- **Sampling strategy:** *which* samples are combined



Samples with the same track role do not overlap

$$\text{trackRole}_i = \text{trackRole}_j \implies \neg \text{OVERLAP}(\text{sample}_i, \text{sample}_j) \quad \forall i \neq j$$

where

$$\text{OVERLAP}(\text{sample}_i, \text{sample}_j) \iff \text{end}_i > \text{start}_j \wedge \text{end}_j > \text{start}_i$$

## Overlapping samples must start at the same time

$$\text{OVERLAP}(\text{sample}_i, \text{sample}_j) \implies \text{start}_i = \text{start}_j \quad \forall i \neq j$$

Rationale: without this constraint overlapping samples might be out-of-sync.

# Constraints

The sum of track role “weights” at each step must be  $\leq$  a threshold

COMULATIVE(*samples*, *demands*, *capacity*)

where *capacity* = 6 and

$$demand_i = \begin{cases} 3 & \text{if } trackRole_i \in \{\text{mainMelody, riff}\} \\ 2 & \text{if } trackRole_i \in \{\text{subMelody, accompaniment}\} \\ 1 & \text{if } trackRole_i \in \{\text{bass, pad}\} \end{cases}$$

Rationale: we want “heavier” track roles to overlap with “lighter” ones.

# Objective

## Minimize makespan

minimize  $\max\{end_i\}$

Rationale: this objective encourages the solver to produce more “compact” solutions — without it, a trivial solution would be to arrange the samples sequentially.

The program is implemented with the CP-SAT interface of the Google OR-Tools framework.

# Sampling strategy

- All selected samples must **share** the **input metadata**.
- ComMU is queried for **one sample for each track role**.
- If there is no sample with input metadata + that track role, one of the available track roles is chosen at random and a sample with the new track role is picked.
- When 6 samples (as many as there are track roles) have been selected, the sampling process is completed.

Of these 6 samples, 3 **will be repeated** throughout the composition: this is achieved by defining an *OptionalIntervalVar* for each sample.

## Extension: Sample Generation

---

# Pros and cons

- MidiComb can also work with samples **generated** by Hyun et al.'s trained Transformer model.
- The **CP model does not change**, only the sample retrieving process does.
- Although sample generation leads to more **diverse** musical compositions, their **quality** is not always on par with their ComMU counterparts.
- A direction for future research might be going beyond the Transformer baseline and devise architectures **tailored to music structure**, instead of relying on preexisting methodologies developed for natural language.



*`frallebini.github.io/midicomb-demo`*